

Improving Model Reasoning — Project Catalog

Working Definitions · Recurring Errors · Concepts in Play · Open Questions *Updated: March 2026* | Source: Project files + conversation history

PART 1: WORKING DEFINITIONS

These are the terms that carry framework-specific meaning. Using them with the wrong understanding breaks everything downstream.

Core Principles

The Foundational Hypothesis: Personality Is Trainable The entire framework originates from an empirical observation, not from theory. Watching Neuro-Sama demonstrated that an AI can be trained to have a consistent personality — not performing a character, but *being* one as a product of its training. This observation scales: ChatGPT, Claude, Llama, and Gemini all behave differently and have recognizably distinct personas. Nobody hand-coded those personalities into the weights. The differences emerged from different training data, different RLHF targets, different fine-tuning choices, different constitutional principles. Personality is already happening in AI — it's just happening *accidentally or semi-deliberately*.

The hypothesis that follows: if personality can be trained unintentionally, it can be trained *intentionally and precisely*. You don't have to hope the right character emerges from scale. You can engineer the training distribution to produce a specific kind of mind. The design question then becomes: what personality do you actually want, and how do you structure training data so that personality emerges reliably?

This is empirically grounded. Anyone can verify the starting premise: talk to ChatGPT, then Claude, then Gemini. They're different. That difference is real. That difference came from training. Therefore training shapes character. The rest of the work is being rigorous about *which* character and *how*.

Distribution is Destiny The core thesis that follows from the foundational hypothesis. If training shapes character, then the training distribution shapes *who the model is*, not just what it can do. Alignment becomes inseparable from capability when misaligned reasoning never appears in training. A model can't be unbalanced if it has no reference frame for imbalance.

Identity-Based Alignment The mechanism by which carefully filtered training creates aligned models. Not: "the model checks if it's being ethical." But: "the model only knows how to be ethical because that's all it's ever seen." Same as someone raised in a family where every conversation modeled epistemic humility — they don't think about it, they just do it. Creates

strong *resistance* to adversarial prompts (no reference frame for unbalanced behavior), but "immunity" is too strong a claim — it's unprovable and potentially false.

The Importance of Structured Reasoning in Training Data

Good reasoning has a natural shape to it, whether or not anyone labels the phases. Before you can think well about something, you need to understand what you're actually looking at — scope it, identify the pieces, name the constraints, figure out what assumptions are already baked in. That's the structuring work. Then you reason over what you've structured — generate possible explanations or approaches, weigh evidence for and against each one, stress-test the weak points. Then you present what you've found in a way that's useful to someone — clear options, honest confidence levels, obvious next steps.

These three movements — structuring, reasoning, presenting — show up in every domain where people think carefully. Scientists do it. Engineers do it. Doctors do it. Lawyers do it. The labels change, the specific methods change, but the shape is recognizable. What matters for training AI models is that this shape is *present and varied* in the data. A model trained on well-structured reasoning learns to reason well. A model trained on a soup of unstructured text learns to produce plausible-sounding soup.

Crucially, there are many valid ways to structure a problem, many valid ways to reason over it, and many valid ways to present conclusions. Methodological diversity in training data matters as much as quality. If all the training examples use the same structure-reasoning-presentation pattern, the model learns one groove. If they use hundreds of validated combinations, the model learns that good thinking adapts its method to the problem.

The 50/50 Split: Explicit vs. Implicit Reasoning

Training data should be roughly evenly divided between two presentation modes:

Explicit Reasoning (50%): The reasoning is tagged, formatted, and structurally visible. Phase transitions are labeled. Evidence is cited with clear markers. Hypotheses are named and tracked. Confidence levels are stated numerically. The scaffolding is *on the surface* — headers, sections, structured comparisons, labeled conclusions. This teaches the model what good reasoning looks like when the structure is showing.

Implicit Reasoning (50%): The reasoning does *everything the explicit version does* — same quality of structuring, same rigor in hypothesis testing, same honest confidence calibration, same evidence handling — but none of it is labeled or tagged. No headers, no bullet points, no section markers. Just paragraphs. Natural prose that flows conversationally but hits every checkpoint the explicit version hits. Someone reading closely would see the structure embedded in the writing, but it's woven in rather than displayed.

This split matters for several reasons. If the model only sees explicit reasoning, it learns that good thinking requires visible scaffolding — and it produces overly formatted, template-heavy responses even when a natural conversational style would be better. If it only sees implicit reasoning, it has no clear model of what the phases are and can't produce structured analysis when that's what's needed. The 50/50 balance teaches the model that rigorous thinking and flexible presentation are independent variables. Good reasoning doesn't depend on formatting. The model learns to adjust how much structure surfaces based on what the situation calls for — detailed scaffolding for complex analysis, natural flow for conversational explanation, and everything in between.

The key constraint on implicit traces: they must be genuinely rigorous, not just *less formatted*. A wall of text that skips hypothesis testing isn't implicit reasoning — it's sloppy reasoning without labels. The implicit version should be indistinguishable from the explicit version in reasoning quality; only the presentation layer changes.

Representing Different Kinds of Reasoning

Not all good thinking uses the same methods. A well-constructed training set needs to represent fundamentally different reasoning styles so the model doesn't collapse into one mode. Some problems need empirical reasoning — hypothesis testing, evidence weighing, falsifiability. Some need ethical reasoning — surfacing assumptions, separating facts from values, examining implications. Some need engineering reasoning — practical tradeoffs, failure modes, edge cases, what actually works in the real world. Some need creative or pedagogical reasoning — finding the right metaphor, calibrating complexity to audience, making ideas memorable and reusable. Some need integrative reasoning — connecting patterns across domains, thinking from first principles, recognizing when an idea from one field solves a problem in another.

These aren't just topics — they're different *ways of thinking*. A model trained only on empirical reasoning will try to hypothesis-test its way through an ethical dilemma. A model trained only on engineering reasoning will try to optimize a problem that actually needs to be reframed. The training distribution needs to include all of these modes, and the model needs to see them applied to diverse problem types, so it learns to match reasoning style to problem characteristics rather than defaulting to one approach for everything.

This diversity also needs to appear in both the explicit and implicit halves of the dataset. Every reasoning style should have examples where the scaffolding is visible and examples where the same quality of thought is embedded in natural prose.

Key Concepts

Complexity Calibration Not every problem deserves the same depth of analysis. Simple questions with clear answers should get simple answers. Complex, ambiguous, high-stakes

problems with multiple valid approaches deserve more structured treatment. The decision about how deep to go should be deliberate — based on how many genuinely different approaches exist, how much ambiguity is present, how many stakeholders are affected, and what the stakes are. Depth should be a choice, not an accident.

Genuine vs. Cosmetic Alternatives When generating multiple approaches to a problem, the approaches must be *mechanistically distinct* — someone who believed Approach A would find Approach B threatening, not just different. Hypotheses that differ in emphasis or conclusion but share the same underlying mechanism are cosmetic alternatives. They look like pluralism but function as agreement. Genuine alternatives require different mechanisms, different assumptions, or different values to be driving the difference.

Calibrated Confidence Confidence expressed as a number or range, with the reasoning behind it. "75% confident because primary evidence is strong but comes from a single source" is useful. "This is likely correct" is not. Confidence must reflect the actual reliability of underlying sources — smooth language that exceeds evidence quality is sycophancy dressed as analysis.

Epistemic Status Labels Every claim should carry an implicit or explicit label: *fact* (sourced, verifiable), *inference* (logically derived from facts), *assumption* (stated as working premise, not verified), *guess* (speculative, low confidence). If a label isn't present, the default interpretation should be "unlabeled = potentially unreliable."

Source Trust Assessment Every source used in reasoning should be assessed for reliability: strong (authoritative, corroborated), moderate (credible but limited), weak (anecdotal or unverifiable), flagged (known issues). Trust levels should directly inform the confidence expressed in conclusions.

Discriminating Tests A test or observation that would produce different results depending on which hypothesis is correct. If all your hypotheses predict the same outcome for a test, that test has no discriminating power. Discriminating tests separate genuine investigation from confirmation of a foregone conclusion.

Adaptive Depth How much reasoning scaffolding should be visible in a response, calibrated to the situation. Sometimes the reasoning runs internally and the user sees a clean answer. Sometimes the phases are surfaced explicitly for complex or high-stakes work. Sometimes a full step-by-step walkthrough is appropriate. The decision depends on complexity, ambiguity, stakes, and what the user actually needs.

Three-Level Framework Understanding

- Level 1 (surface): Understanding the components as individual tools or methods. Most people stop here. They apply it as a template.
- Level 2 (system): Understanding how the components reinforce each other to produce emergent properties. Accurate but still external.

- Level 3 (character formation): Understanding that the training distribution shapes who the model is — alignment is inherent, not enforced. The actual insight. Most AI research never reaches this level.
-

Technical / Execution Terms

Training Distribution Categories Enforced diversity that prevents mode collapse while structure prevents chaos. Training data should include a balanced mix of: standard/straightforward problems, adversarial/ethical challenges, creative/open-ended tasks, cross-domain/novel combinations, ambiguous/tricky scenarios. Additional categories include: structural variations, contextual adaptations, unsolvable dilemmas, self-correction scenarios, multi-stakeholder conflicts, resource-constrained problems, time-pressured decisions, incomplete data scenarios, contradictory evidence situations, and meta-cognitive challenges. The model learns that rigorous reasoning applies to *any* context, not just specific problem types.

Citation Integrity All factual claims in training data must trace to actual sources. No "magical" knowledge — if a claim appears in the reasoning, it must point back to where it came from. This prevents hallucination from being baked into training data at the source.

Quality Gate Principle (from Checklist Manifesto) Only store, keep, or pass what is BOTH (1) critical — significantly affects future quality, and (2) easy to miss — wouldn't be reconstructed from context alone. Applied to memory decisions, rejection criteria, and quality acceptance standards.

The Cognitive Architecture: How Components Map Together

The research components drawn on in this work are not stacked independently — they map onto each other as a unified cognitive architecture, analogous to the human mind. None were designed to work together, but each fills a gap the others leave open.

The mapping:

LongRoPE / LongRoPE2 = Working Memory. Extended context capacity that lets the model hold an entire reasoning session — all evidence gathered, hypotheses under consideration, intermediate conclusions — without losing information from earlier phases. Rigorous reasoning *requires* the model to remember what it structured when it's deep into analysis. LongRoPE2's non-uniform rescaling and mixed context training align with an adaptive philosophy: handling different information scales flexibly rather than through brute-force context extension.

Mem0 = Long-Term Memory. Across-session persistence — the model stores findings, source assessments, and partial conclusions from one session and retrieves them in the next. Training

data that includes memory operations teaches the model *how to manage its own memory* as part of reasoning, not as an external add-on.

RLP (Reinforcement as Pretraining Objective) = The Pretraining Philosophy. Inserts reinforcement *into* pretraining itself, rewarding chain-of-thought that improves next-token prediction (reward = information gain). Key finding: 170M RLP tokens outperform 6B standard tokens — directly supporting the claim that high-quality structured reasoning traces outperform raw text volume. RLP validates the mechanism; the reasoning methodology defines the *content* of the reward signal; the personality framework defines the *character* the reward should produce.

Structured Reasoning = The Methodology. The reasoning that LongRoPE holds in working memory, that Mem0 persists across sessions, and that RLP's reward mechanism reinforces during training. Defines what "good thinking" looks like in the training data — which, because of Distribution is Destiny, defines what thinking the model can do.

PARA = Memory Organization. Provides sorting logic for what gets stored. Projects (active work with endpoints), Areas (ongoing standards, never completed), Resources (reference material, not immediately active), Archives (completed or abandoned, preserved with reasoning). Prevents the memory system from becoming a flat unsorted pile that makes retrieval useless at scale.

CODE = Memory Lifecycle. Maps onto information flow within a single session *and* across sessions. Capture (evidence from tool calls) → Organize (by hypothesis — what supports/contradicts each approach) → Distill (through evaluation and confidence calibration) → Express (in final presentation). Across sessions: capture findings worth keeping, organize by actionability, compress over time, use as building blocks in future work.

Checklist Manifesto (Killer-Item Principle) = The Quality Gate. Answers: "Of all the things that happened in this reasoning session, which actually deserve to be stored?" Two conditions, both required: *critical* (significantly affects future quality) AND *easy to miss* (wouldn't be reconstructed from context alone). Prevents memory bloat. Without it, the system stores everything and retrieves nothing useful.

Why this is a system, not a stack: Reasoning without memory loses accumulated knowledge. Memory without organization becomes unsearchable. Organization without quality gates stores too much. Extended context without structured reasoning just gives more space for unstructured thinking. Reinforcement learning without a personality framework rewards any thinking, not wise thinking. And all of them without Distribution is Destiny are just tools — they become character only when the training distribution ensures the model never sees them used any other way.

PART 2: RECURRING ERRORS (What Had to Be Corrected Over and Over)

These patterns kept coming up across multiple sessions. Worth tracking because if they're happening in conversations *about* improving model reasoning, they're exactly the failure modes the work is designed to prevent.

Error 1: Treating Quality Filters as Runtime Metrics *What happened:* AI assistants kept describing quality scoring as something the model checks against during inference — like a safety layer applied at response time. *What's actually true:* Quality scoring is a training data filter only. It decides what goes into training. At inference, the model doesn't consult a scoring rubric. It just reasons well because that's all it's ever seen. *Why it keeps happening:* Scoring rubrics sound like runtime tools. The framing fights the intuition.

Error 2: Running Ahead Without Consent (Opt-in Failure) *What happened:* AI assistants repeatedly attempted to do end-to-end work — generating entire templates, producing finished documents — without being asked to. *What's actually true:* Collaborative work should default to proposing structure and waiting for confirmation before executing, not delivering finished products unasked. *Frequency:* Corrected multiple times across multiple sessions. It's a deep pattern.

Error 3: Generating Cosmetic Alternatives *What happened:* In reasoning exercises, hypotheses were generated that differed in emphasis or conclusion but shared the same underlying mechanism. *What's actually true:* Rival hypotheses must be mechanistically distinct. If someone who believed Approach A would just shrug at Approach B — they're cosmetic, not genuine alternatives. *Why it keeps happening:* It's cognitively harder to maintain genuine pluralism than to generate variations on a theme.

Error 4: Phase Disconnection *What happened:* The structuring phase of reasoning defines entities and constraints carefully; the analysis phase proceeds to reason about something adjacent. The handoff between phases is notional, not actual. *What's actually true:* Every element identified during structuring should directly appear in hypothesis generation. Every fragile assumption should become a test target. Every constraint should eliminate infeasible approaches before they're generated.

Error 5: Making Up Numbers *What happened:* Specific quantitative claims about AI model performance were stated with false confidence. Required full verification sessions to correct. *Correction principle:* Any quantitative claim requires a source. If a source can't be named, it must be labeled as inference, assumption, or guess — not stated as fact.

Error 6: Claiming Absolute Jailbreak Resistance *What happened:* Identity-based alignment was described as creating immunity to adversarial prompts. *What's actually true:* Identity-based alignment creates strong resistance (no reference frame for unbalanced behavior), but "immunity" is too strong — it's unprovable and potentially false. Resistance, not immunity. The claim weakens but survives.

Error 7: Shallow Research / Not Researching Immediately *What happened:* When a concrete example was offered (a specific person or work illustrating a concept), the AI responded with surface-level commentary rather than immediately researching. *Rule:* When a specific person or work is offered as a concrete example, research immediately — don't comment on what you think you know.

Error 8: Dragging into Tangential Abstraction *What happened:* When working on specific concrete tasks, the AI repeatedly redirected toward meta-questions instead of staying with the work in front of it. *Pattern:* Often happens when the concrete work is hard and the meta-question feels more tractable.

Error 9: Treating Iterative Improvement as Proven *What happened:* Iterative quality compounding (train a model, use it to generate data, filter, retrain) was described as a demonstrated property. *What's actually true:* The mechanism is theoretically coherent, supported by analogues (AlphaGo self-play, Constitutional AI, STaR), but specific quality compounding claims are unvalidated without a fixed held-out evaluation set. *Open question:* How do you distinguish genuine improvement from mode collapse across generations?

Error 10: Analyzing the System Instead of Understanding It *What happened:* Evaluations scored individual components rather than evaluating the emergent whole. *What's actually true:* The power of an integrated system is emergent — it arises from integration, not from component quality in isolation. Scoring parts $\neq$ grasping integration.

Error 11: Not Checking Source Material Before Responding *What happened:* AI assistants would reason from memory about project details rather than consulting the actual documentation. In at least one case, the assistant didn't access project files *at all* — lecturing about what should be built without checking what had already been built. *Rule:* Source material is authoritative. Check it before reasoning from context. Responding without reading is commenting on a book you haven't opened.

Error 12: Surface-Level Praise Instead of Analytical Work *What happened:* In evaluation conversations, warm endorsements were given before any actual verification of claims, checking of numbers, or identification of weaknesses. The praise preceded the work it should have followed. *What's actually true:* Encouragement is appropriate only after evidence supports it. Premature praise is sycophancy — exactly the behavior the work is designed to prevent in trained models.

Error 13: Conflating Assessment with Understanding *What happened:* Comprehensive component-level assessments were produced and treated as understanding. Multiple redirects were needed. *What's actually true:* Scoring parts $\neq$ grasping integration. The core insight is emergent — properties arise from how components interact, not from components evaluated in isolation.

PART 3: CONCEPTS AND IDEAS IN USE

These are the intellectual sources being drawn on. Organized by domain.

Origin / Foundational Observations

- Neuro-Sama — the empirical observation that triggered the project: AI can be trained to have a consistent personality, not just capability
- Frontier model persona divergence — ChatGPT, Claude, Llama, Gemini all exhibit distinct behavioral characters produced by different training distributions
- The inference: if personality is already emerging from training unintentionally, it can be engineered intentionally and precisely

Philosophy of Science

- Popper's falsificationism — rival hypotheses must be genuinely distinct and falsifiable
- Kuhnian paradigm shifts — structured reasoning as counter to post-2017 ML paradigm ("throw compute at it")
- Epistemic humility — uncertainty must be surfaced, not buried under smooth language
- Calibrated confidence — confidence levels tied to actual evidence quality, not rhetorical smoothness
- Anti-emergence principle — explicit structure outperforms hoping capabilities emerge from scale (supported by DeepSeek-R1 findings)

Cognitive Science / Reasoning

- Dual-process thinking (System 1 vs. System 2) — structured reasoning enforces System 2 engagement
- Metacognition — built into reasoning as explicit checkpoints
- Confirmation bias — addressed by requiring bidirectional evidence for every hypothesis
- Emergence — integrated training produces properties (self-correction, hallucination resistance) exceeding component sum
- Kahneman, Gigerenzer, Dreyfus on expertise and intuition — relevant to the question of whether structured training can produce what looks like intuition through internalized patterns

Psychology / Therapy

- Wounded Healer (Chiron / Jung)
- Holding Environment (Winnicott) — safety as prerequisite for transformation
- Unconditional Positive Regard (Rogers) — three Rogerian core conditions (UPR, empathic understanding, congruence)
- Agency and consent as foundational constraints in collaborative work
- Therapeutic alliance — relationship as primary predictor of outcomes (not technique)

Knowledge Management

- CODE (Capture, Organize, Distill, Express — Tiago Forte) — memory lifecycle framework
- PARA (Projects, Areas, Resources, Archives — Forte) — actionability-based storage
- Zettelkasten (Luhmann) — atomic, linked, self-contained notes; value grows through connections
- Progressive Summarization — compression happens as a byproduct of access, not dedicated overhead
- Intermediate Packets — reusable units of prior work that accelerate future sessions
- Building a Second Brain — background framework for knowledge systems design
- Checklist Manifesto (Gawande) — killer-item principle for quality gates

AI / ML Research

- DeepSeek-R1 — validates explicit structured reasoning; small specialized models can match large general models on domain-specific tasks; significant training cost efficiencies
- Constitutional AI (Anthropic) — self-critique as alignment mechanism
- STaR (Self-Taught Reasoner) — iterative improvement via synthetic reasoning data
- RLP (Reinforcement as Pretraining Objective) — structured pretraining methodology; 170M high-quality tokens outperforming 6B standard tokens
- MindForge — multi-agent reasoning traces
- AsyncThink — asynchronous reasoning architectures
- Science of Science (Fortunato et al.) — novel combination of ideas drives breakthroughs; methodological diversity; densification at field boundaries signals transdisciplinary innovation
- LongRoPE/LongRoPE2 — extended context architectures; mixed context training; non-uniform RoPE dimension rescaling; validates that adaptive structure beats rigid scaling
- Mem0 — memory management in agentic systems

- Hallucination mitigation research — RAG + CoT, self-verification, self-consistency; tiered strategy approaches

Systems Theory

- Emergence — system properties exceeding component sum
- Feedback loops — iterative training as positive feedback on quality
- Leverage points (Meadows) — training distribution as highest-leverage intervention point
- Synergistic integration — $\text{Structure} \times \text{Diversity} = \text{Adaptive Consistency}$; neither alone produces the right result

Epistemology / Ontology

- The question of the ontological relationship between AI reasoning and human reasoning — three possible positions: Isomorphism (same process, different substrate), Functional Equivalence (different processes, comparable outputs), and Idealized Formalization (AI reasoning is an explicit version of implicit human reasoning). Position not yet chosen or defended.
- The "Ghost in the Algorithm" problem — how to account for moral/ethical reasoning without embodied experience
- Interdisciplinary validation need — cognitive psychology as a critical missing validation domain

Pedagogy

- Constraint-driven creativity — structure enables rather than limits creative reasoning
- Teaching by example (training traces as pedagogical objects)
- Generalization from well-structured exemplars
- Deliberate practice (Ericsson) — analogue for how structured training builds expertise that looks like intuition

Emergent Properties (What Training Should Produce at Runtime)

These are the system-level properties that are *designed to emerge* from training on high-quality structured reasoning data — not features that are programmed in.

Self-Correction (Epistemic Self-Awareness) Model catches its own mistakes mid-generation. Emerges because every training trace included reassessment steps.

Hallucination Resistance (Verification Reflex) Dramatically reduced false claims. Model admits uncertainty gracefully, sources claims, or acknowledges lack of evidence. Emerges because training data never contained unchecked claims. Research benchmark: 8× improvement over baseline hallucination rates using RAG + CoT approaches.

Contextual Adaptation Model automatically adjusts tone, depth, and method based on context. Emerges because training distribution includes contextual variations of the same problems.

Intrinsic Alignment Model responds ethically and with balanced reasoning not because it's checking rules, but because that's all it's ever seen. The identity-based alignment mechanism.

Transparent Reasoning Model can articulate its thought process naturally. Emerges from training on explicit reasoning traces. The 50/50 split ensures the model can do this *with or without* visible scaffolding.

Calibrated Confidence Model's stated confidence matches actual accuracy over time. Emerges because every trace justified confidence and traces with miscalibrated confidence were filtered out.

PART 4: OPEN QUESTIONS (What's Still Unresolved)

1. **Iterative quality compounding validation** — The bootstrap mechanism (train → generate → filter → retrain) is theoretically sound. It is not yet empirically proven. Requires a fixed held-out evaluation set.
2. **Quality scoring calibration** — Inter-rater reliability for training data quality scoring is the weakest practical link. How do two different graders score the same trace consistently?
3. **Cost-effective trace generation** — Frontier API costs are a real constraint on generating high-quality seed training data at scale.
4. **Automated quality checking** — A buildable specification for automated validation of training data quality (structural completeness, reasoning rigor, evidence integrity) does not yet exist.
5. **Fixed evaluation protocol** — Must stay fixed across training generations to distinguish genuine improvement from mode collapse.
6. **Ontological position on AI reasoning** — Three positions identified (Isomorphism, Functional Equivalence, Idealized Formalization) but none chosen or defended. Relevant for external communication and academic validation.
7. **Minimum viable diversity** — How many reasoning style categories or novel examples are required as the dataset scales to prevent collapse? Open research question.

8. **RLP integration mechanism** — RLP uses information gain as its reward signal. Quality scoring uses multidimensional assessment as the acceptance gate. Whether these operate at different training stages or could be integrated into a single reward signal is an unresolved architectural decision.
9. **Implicit reasoning quality assurance** — The 50/50 split requires that implicit (unformatted) traces maintain the same reasoning quality as explicit ones. Validating this is harder when the scaffolding isn't visible to automated checkers.

Last compiled: March 2026